

CS338 Assignment 4

Kuo Wang

April 2022

When I paste the address and click “go” on the browser, the browser connects to a DNS server first to exchange the cs338.jeffondich.com domain name for an IP address: 45.79.89.123. Both an A and AAAA record requests are made, and both return the same result. AAAA records, for additional knowledge, is able to resolve IPv6 addresses as a newer version and eventual replacement of A records.

```
1  0.000000000  172.16.34.128 172.16.34.2  DNS    80      Standard query 0x39c9 A cs338.
   jeffondich.com
2  0.000071937  172.16.34.128 172.16.34.2  DNS    80      Standard query 0x9ec5 AAAA cs338.
   jeffondich.com
3  0.000746153  172.16.34.2   172.16.34.128 DNS    96      Standard query response 0x39c9 A
   cs338.jeffondich.com A 45.79.89.123
4  0.001683400  172.16.34.2   172.16.34.128 DNS    80      Standard query response 0x9ec5
   AAAA cs338.jeffondich.com
```

Now that the browser has the IP address, it initiates the standard DNS handshake with the given IP address. The first SYN TCP packet is sent from client port 50880 to server port 80.

```
5  0.001907226  172.16.34.128 45.79.89.123 TCP    74      50880    80 [SYN] Seq=0 Win=64240
   Len=0 MSS=1460 SACK_PERM=1 TSval=2223634628 TSecr=0 WS=128
```

Almost simultaneously, another TCP SYN packet is sent from port 50882 to port 80. Apparently the web browser sends multiple packets through different ports to increase load speed with parallelism.

```
6  0.007191050  172.16.34.128 45.79.89.123 TCP    74      50882    80 [SYN] Seq=0 Win=64240
   Len=0 MSS=1460 SACK_PERM=1 TSval=2223634634 TSecr=0 WS=128
```

Continuing on, the TCP handshake through both ports carry on. For the transactions between client port 50880 to server port 80, SYN is met with SYN,ACK from the server, then the client sends an ACK as response. A connection is established here.

```
7  0.092932245  45.79.89.123 172.16.34.128 TCP    60      80      50880 [SYN, ACK] Seq=0 Ack=1
   Win=64240 Len=0 MSS=1460
8  0.092964556  172.16.34.128 45.79.89.123 TCP    54      50880    80 [ACK] Seq=1 Ack=1 Win
   =64240 Len=0
```

The same goes for 50882 for initial connection with TCP.

```

6  0.007191050  172.16.34.128 45.79.89.123 TCP    74      50882    80 [SYN] Seq=0 Win=64240
    Len=0 MSS=1460 SACK_PERM=1 TSval=2223634634 TSecr=0 WS=128
11 0.100864836  45.79.89.123 172.16.34.128 TCP    60      80      50882 [SYN, ACK] Seq=0 Ack=1
    Win=64240 Len=0 MSS=1460
12 0.100890806  172.16.34.128 45.79.89.123 TCP    54      50882    80 [ACK] Seq=1 Ack=1 Win
    =64240 Len=0

```

Then, through 50880, the client posts an HTTP GET request to the server for the path `/basicauth/`. The server responds with a TCP ACK packet signaling that it received the packet (HTTP uses TCP as underlying protocol).

```

9  0.093187171  172.16.34.128 45.79.89.123 HTTP    395     GET /basicauth/ HTTP/1.1
10 0.093278765  45.79.89.123 172.16.34.128 TCP    60      80      50880 [ACK] Seq=1 Ack=342
    Win=64240 Len=0

```

Then, the server sends a HTTP protocol packet with status code 401 in its header to tell the browser that it is not authorized to get a response from the server.

```

13 0.245854333  45.79.89.123 172.16.34.128 HTTP    457     HTTP/1.1 401 Unauthorized (text/
    html)

```

The browser sees this, sends an ACK to signal the server, and asks me for credentials.

```

14 0.245885012  172.16.34.128 45.79.89.123 TCP    54      50880    80 [ACK] Seq=342 Ack=404
    Win=63837 Len=0

```

I key in the right credentials, and click submit. This triggers the browser to submit another HTTP GET request with an Authorization header...

```

21 3.564746561  172.16.34.128 45.79.89.123 HTTP    438     GET /basicauth/ HTTP/1.1

```

...which includes the credentials in base64 text. "Basic" lets any application involved in the process know that the HTTP authorization scheme is basic: just a pair of username/password. Then, "Y3MzMzg6cGFzc3dvcmQ=" is the combination itself in Base64 format, which can be decoded by anyone with a simple tool into the string "cs338:password" in the format of `username:password`. The string is used by the server to verify identity and grant access. Finally, "\r\n" just signals the end of the header by a way outlined by RFC standards.

```

Authorization: Basic Y3MzMzg6cGFzc3dvcmQ=\r\n
Credentials: cs338:password
\r\n

```

Base64 encoding can be decoded into plain string by anyone; there is no secrets involved in the decoding process. This is a problem, because essentially the lack of end-to-end encryption means that whoever gets the package can know the credentials; therefore, the system is prone to packet sniffing exploits.

Anyway, the server now has a GET request from the client with the right credentials. The server sends to the client an ACK so it knows that the packet was received. Afterwards, as the server verified the credentials, my browser and I are let in through the invitation of an HTTP packet with status code 200 (OK) along with the data I requested. Of course, the browser reciprocated the HTTP packet with another ACK packet.

```

21 3.564746561 172.16.34.128 45.79.89.123 HTTP 438 GET /basicauth/ HTTP/1.1
22 3.565075098 45.79.89.123 172.16.34.128 TCP 60 80 50880 [ACK] Seq=404 Ack=726
    Win=64240 Len=0
23 3.650702440 45.79.89.123 172.16.34.128 HTTP 458 HTTP/1.1 200 OK (text/html)
24 3.650721025 172.16.34.128 45.79.89.123 TCP 54 50880 80 [ACK] Seq=726 Ack=808
    Win=63837 Len=0

```

The data is an HTML page that links to three other sites through hyperlinks. This is found in the HTTP data packet as Line-based text data.

Simultaneously, 50882 is also working in parallel. As the webpage is loaded, the favicon is requested by the browser through the 50882 port, which was kept alive. The favicon must not have been set up to require authorization, because 50882 neither asked for credential nor contain authorization header. Since there isn't a favicon, a RST is sent from the client after the HTTP 404 response.

```

25 3.738839869 172.16.34.128 45.79.89.123 HTTP 355 GET /favicon.ico HTTP/1.1
27 3.739008701 45.79.89.123 172.16.34.128 TCP 60 80 50882 [ACK] Seq=1 Ack=302
    Win=64240 Len=0
29 3.932877042 45.79.89.123 172.16.34.128 HTTP 383 HTTP/1.1 404 Not Found (text/html)
30 3.932905446 172.16.34.128 45.79.89.123 TCP 54 50882 80 [RST] Seq=303 Win=0
    Len=0

```

So now, I am at the web page which is successfully loaded. The client and the server regularly exchange TCP keep-alive packets to ensure that the connection exists and further actions are possible.

```

35 13.757264492 172.16.34.128 45.79.89.123 TCP 54 [TCP Keep-Alive] 50880 80 [ACK]
    Seq=725 Ack=808 Win=63837 Len=0
36 14.781214430 172.16.34.128 45.79.89.123 TCP 54 [TCP Keep-Alive] 50880 80 [ACK]
    Seq=725 Ack=808 Win=63837 Len=0
37 14.781522898 45.79.89.123 172.16.34.128 TCP 60 [TCP Keep-Alive ACK] 80 50880 [
    ACK] Seq=808 Ack=726 Win=64240 Len=0

```

When I click on a link, my browser sends to the server through port 50880 (client) to 80 (server) a new HTTP GET request. The GET request also contains the same authorization header as before, so access to the destination file is also restricted. As the first page finished loading, the browser reset the connection from its port 50880. I closed the tab, and that might have triggered the RST.

```

42 32.805532807 172.16.34.128 45.79.89.123 HTTP 499 GET /basicauth/amateurs.txt HTTP
    /1.1
44 32.805791820 45.79.89.123 172.16.34.128 TCP 60 80 50880 [ACK] Seq=808 Ack=1171
    Win=64240 Len=0
45 32.805791941 45.79.89.123 172.16.34.128 TCP 60 80 50880 [ACK] Seq=808 Ack=1172
    Win=64239 Len=0
46 32.911529940 45.79.89.123 172.16.34.128 HTTP 375 HTTP/1.1 200 OK (text/plain)
47 32.911564155 172.16.34.128 45.79.89.123 TCP 54 50880 80 [RST] Seq=1172 Win=0
    Len=0

```

Accessing all of amateurs.txt, dancing.txt, and armed-guards.txt follow the same procedure, just with a new connection established from a different client port to the same server port. The new connection itself went through the regular TCP handshake procedures to establish.

Afterwards, the processes are always in the order of client sending HTTP GET to the server with path to the desired file, server sends back ACK, server sends back HTTP response with code 200 and desired file, followed by client ACK. The connection was never reset.

As expected, TCP keep alive transactions are present and sporadic.

```

62 37.796176669 172.16.34.128 45.79.89.123 HTTP 503 GET /basicauth/armed-guards.txt
    HTTP/1.1
63 37.796432336 45.79.89.123 172.16.34.128 TCP 60 80 50922 [ACK] Seq=1129 Ack
    =1620 Win=64240 Len=0
64 37.910365110 45.79.89.123 172.16.34.128 HTTP 462 HTTP/1.1 200 OK (text/plain)
65 37.910383646 172.16.34.128 45.79.89.123 TCP 54 50922 80 [ACK] Seq=1620 Ack
    =1537 Win=63837 Len=0
66 38.492003610 172.16.34.128 45.79.89.123 TCP 54 [TCP Keep-Alive] 50924 80 [ACK]
    Seq=301 Ack=330 Win=63911 Len=0
67 39.159604036 172.16.34.128 45.79.89.123 HTTP 498 GET /basicauth/dancing.txt HTTP
    /1.1
68 39.159769722 45.79.89.123 172.16.34.128 TCP 60 80 50922 [ACK] Seq=1537 Ack
    =2064 Win=64240 Len=0
69 39.343699813 45.79.89.123 172.16.34.128 HTTP 528 HTTP/1.1 200 OK (text/plain)
70 39.343718959 172.16.34.128 45.79.89.123 TCP 54 50922 80 [ACK] Seq=2064 Ack
    =2011 Win=63837 Len=0

```

Apparantly, all three files have restricted access since the GET requests have authorization headers, and attempts at accessing them with no authorization proved futile.